

# Backup slides

Md. Sakif Khan    0421052099

- 2    Budget configuration and enforcement sites
- 3    Which budgets actually bind
- 4    The full failure census
- 5    Hedging behaviour, all five systems

## Backup: budget configuration

| Parameter        | Value | Enforced in                        |
|------------------|-------|------------------------------------|
| max_depth        | 4     | both routers                       |
| beam_width       | 3     | explorer (slice)                   |
| max_anchors      | 5     | explorer <i>and</i> evaluator      |
| max_backtracks   | 3     | post-evaluation router             |
| max_llm_calls    | 25    | LLM client                         |
| max_verify_iters | 2     | post-verification router, verifier |
| max_seconds      | 300   | both routers                       |


Termination does not depend on model behaviour: every cycle passes through a router that checks a monotone counter.

## Backup: which budgets actually bind


| Budget               | WebQSP      | CWQ         | Pooled      |
|----------------------|-------------|-------------|-------------|
| Depth cap            | 16.5%       | 34.8%       | 25.6%       |
| Backtrack cap        | 3.2%        | 9.2%        | 6.2%        |
| Verify-iteration cap | 1.0%        | 2.8%        | 1.9%        |
| <b>Call cap</b>      | <b>0.0%</b> | <b>0.0%</b> | <b>0.0%</b> |

Percentage of questions on which each budget refused a request. AGR never reaches the 25-call cap — which is what makes the comparison against a clipped Think-on-Graph meaningful.

## Backup: the full failure census



## Backup: accuracy against cost



On *tokens* Think-on-Graph is **not** dominated — it buys lower accuracy at a lower token price. On *calls*, which the budget meters, it is.

The frontier depends on which cost you charge.

## Backup: the benchmark was wrong 57 times

The same cross-system agreement is also a **detector**: consensus flagged 105 questions, adjudication confirmed 41, and the gap is the attractor above rather than label noise.

- 41 excluded before the census
- 17 found inside the census
- 1 counted in both

⇒ 57 **distinct questions**

### Why report it

Reported defect rates for these two standard benchmarks are rare and usually anecdotal. This is a counted rate with a documented adjudication procedure and per-question provenance.

**A contribution that outlives this particular system.**