


Every answer passes through a structural verifier.



Measured on 400 held-out questions per benchmark, against four baselines, in one frozen environment

0.755 / 0.522

Hits@1, WebQSP / ComplexWebQuestions
strongest baseline of four: 0.660 / 0.435

0.46 → 0.55 → 0.57

CWQ accuracy across 1, 2 and 3+ hop questions
the only system that ends above where it started

0.0%

of 1,709 asserted entities are absent from the graph
a property of graph navigation, not of the layer

57

questions where the benchmark,
not the system, was at fault